

# MiniSQL个人报告

3220103450 姜雨童

Date: 2024.06.24

# 目录

- 一、完成模块说明
- 二、Catalog Manager模块
  - (一) Catalog Manager概述
  - (二) 模块调用
  - (三) 实现的接口
  - (四) 具体实现
- 三、Planner and executor模块
  - (一) Planner and executor概述
  - (二) 模块调用
  - (三) 实现的接口
  - (四) 具体实现
- 四、模块功能验证
  - (一) Catalog Manager
  - (二) Execute Engine
- 五、代码实现
  - (一) Catalog Manager
  - (二) Planner and executor

## 一、完成模块说明

1. 参与小组讨论，对MiniSQL的实现进行分析
2. 负责模块四（Catalog Manager）和模块五（Executor）的实现

## 二、Catalog Manager模块

### （一）Catalog Manager概述

Catalog Manager 负责管理和维护数据库的所有模式信息，包括：

- 数据库中所有表的定义信息，包括表的名称、表中字段（列）数、主键、定义在该表上的索引。
- 表中每个字段的定义信息，包括字段类型、是否唯一等。
- 数据库中所有索引的定义，包括所属表、索引建立在那个字段上等。

这些模式信息在被创建、修改和删除后还应被持久化到数据库文件中。此外，Catalog Manager还需要为上层的执行器Executor提供公共接口以供执行器获取目录信息并生成执行计划。

### （二）模块调用

- 调用record manager中column, row和schema的序列化和反序列化函数
- 调用disk manager中的bufferpoolmanager，为表和索引分配单独的数据页
- 调用index manager中的BplusTreeIndex，为表中的列创建索引
- （本模块的CatalogManager类为Execute Engine提供函数接口，以实现minisql各模块功能）

### （三）实现的接口

#### 1. catalog.cpp

- `uint32_t CatalogMeta::GetSerializedSize() const`
- `CatalogManager::CatalogManager(BufferPoolManager *buffer_pool_manager, LockManager *lock_manager, LogManager *log_manager, bool init)`

- `dberr_t CatalogManager::CreateTable(const string &table_name, TableSchema *schema, Txn *txn, TableInfo *&table_info)`
- `dberr_t CatalogManager::GetTable(const string &table_name, TableInfo *&table_info)`
- `dberr_t CatalogManager::GetTables(vector<TableInfo *> &tables) const`
- `dberr_t CatalogManager::CreateIndex(const std::string &table_name, const string &index_name, const std::vector<std::string> &index_keys, Txn *txn, IndexInfo *&index_info, const string &index_type)`
- `dberr_t CatalogManager::GetIndex(const std::string &table_name, const std::string &index_name, IndexInfo *&index_info) const`
- `dberr_t CatalogManager::GetTableIndexes(const std::string &table_name, std::vector<IndexInfo *> &indexes) const`
- `dberr_t CatalogManager::DropTable(const string &table_name)`
- `dberr_t CatalogManager::DropIndex(const string &table_name, const string &index_name)`
- `dberr_t CatalogManager::FlushCatalogMetaPage() const`
- `dberr_t CatalogManager::LoadTable(const table_id_t table_id, const page_id_t page_id)`
- `dberr_t CatalogManager::LoadIndex(const index_id_t index_id, const page_id_t page_id)`
- `dberr_t CatalogManager::GetTable(const table_id_t table_id, TableInfo *&table_info)`

## 2. index.cpp

- `uint32_t IndexMetadata::GetSerializedSize() const`（模块内调用）

### 3. index.h

- `void Init(IndexMetadata *meta_data, TableInfo *table_info, BufferPoolManager *buffer_pool_manager)`（模块内调用）

#### （四）具体实现

可以把上述函数接口大致分为以下几类：

- 序列化和反序列化（包括初始化）
- Catalog Manager
- Create Table/Index
- Get Table(s)/Index/TableIndexes
- Load Table/Index
- Drop Table/Index
- Flush CatalogMetaPage

#### 1. 序列化和反序列化

数据库中定义的表和索引在内存中以TableInfo和IndexInfo的形式表现。以IndexInfo为例，它包含了这个索引定义时的元信息meta\_data\_，该索引对应的表信息table\_info\_，该索引的模式信息key\_schema\_和索引操作对象index\_。除元信息meta\_data\_外，其它的信息都是通过反序列化后的元信息生成的。

因此，为了能够将所有表和索引的定义信息持久化到数据库文件并在重启时从数据库文件中恢复，我们需要为表和索引的元信息TableMetadata和IndexMetadata实现序列化和反序列化操作。（为了简便处理，在序列化时我们为每一个表和索引都分配一个单独的数据页用于存储序列化数据。）CatalogMeta的信息将会被序列化到数据库文件的第CATALOG\_META\_PAGE\_ID号数据页中（逻辑意义上），CATALOG\_META\_PAGE\_ID默认值为0。

就具体实现而言，CatalogMeta类的序列化和反序列化函数都已给出，只需要实现几个获取序列化长度的函数以及Index的初始化（Init）。代码实现将统一附在报告最后。

## 2. Catalog Manager

`CatalogManager`能够在数据库实例（`DBStorageEngine`）初次创建时（`init = true`）初始化元数据；并在后续重新打开数据库实例时，从数据库文件中加载所有的表和索引信息，构建`TableInfo`和`IndexInfo`信息置于内存中。

根据数据库的调用情况，可以分为首次创建数据库和重新加载数据库两种情况。

首次调用时，只需要`catalogmanager`进行初始化。而重新加载时，则要通过反序列化得到数据库的元信息。接着通过对`catalogmeta`中信息的遍历，反序列化得到`tablemetadata`和`indexmetadata`，构建对应的`tableinfo`和`indexinfo`并把信息加入`catalogmanager`中，从而加载出完整的`catalogmanager`。

## 3. Create Table/Index

这两个构建函数的实现逻辑类似，这里以`create table`为例。首先需要利用`bufferpoolmanager`获取一个空闲页作为新建表的数据页（对应`page_id`），而调用`catalogmeta`中`GetNextTableId`得到的是新建表的`id`。然后调用初始化函数生成`tableinfo`，并把新建的表存到`table`列表中。最后更新`catalogmetapage`的信息，保证下一次调用时信息是正确的。（注意在实现`create table`的时候要进行深度拷贝，否则会报错。）

## 4. Get Table(s)/Index/TableIndexes

同样以`get table`为例，根据给出的`tablename`（或是`tableid`，这里实现了通过`name`和`id`两个不同标识查找表的函数）在`table`列表`table_`中找到对应`tableinfo`即可。而对于获得多个信息的`get tables/tableindexes`，只需创建向量表存入信息。

## 5. Load Table/Index

以`load table`为例。根据表`id`得到数据页`id`，然后通过反序列化得到`tablemeta`，后续把`tablemeta`中的数据，如`tablename`，`tableinfo`载入到数据库中。

## 6. Drop Table/Index

以`drop table`为例，根据表名，在`catalogmanager`类中`table`名称向量和`table`向量表中删除该表的各类信息。同时要在`catalogmeta`中删除该表，并调用`bufferpoolmanager`类的`deletepage`函数对该表所在的页面进行回收，避免空间上的浪费。

## 7. Flush CatalogMetaPage

把目前catalogmetadata写入数据页，并即时将数据序列化到磁盘中。

### 三、Planner and executor模块

#### (一) Planner and executor概述

本实验主要包括Planner和Executor两部分。

Planner的主要功能是将解释器（Parser）生成的语法树，改写成数据库可以理解的数据结构。在这个过程中，我们会将所有sql语句中的标识符（Identifier）解析成没有歧义的实体，即各种C++的类，并通过Catalog Manager 提供的信息生成执行计划。（这部分没有代码需要实现，只需要了解parser的实现方式，便于后续编写代码。）

Executor遍历查询计划树，将树上的 PlanNode 替换成对应的 Executor，随后调用 Record Manager、Index Manager 和 Catalog Manager 提供的相应接口进行执行，并将执行结果返回给上层模块。

#### (二) 模块调用

- 调用catalog.cpp中CatalogManager类的各函数，以实现minisql各模块的功能

#### (三) 实现的接口

##### 1. execute\_engine.cpp

这些函数均由类内函数ExecuteEngine::Execute（已给出实现）统一调用

- dberr\_t ExecuteEngine::ExecuteCreateTable(pSyntaxNode ast, ExecuteContext \*context)
- dberr\_t ExecuteEngine::ExecuteDropTable(pSyntaxNode ast, ExecuteContext \*context)
- dberr\_t ExecuteEngine::ExecuteShowIndexes(pSyntaxNode ast, ExecuteContext \*context)
- dberr\_t ExecuteEngine::ExecuteCreateIndex(pSyntaxNode ast, ExecuteContext \*context)
- dberr\_t ExecuteEngine::ExecuteDropIndex(pSyntaxNode ast, ExecuteContext \*context)

- dberr\_t ExecuteEngine::ExecuteExecfile(pSyntaxNode ast, ExecuteContext \*context)
- dberr\_t ExecuteEngine::ExecuteQuit(pSyntaxNode ast, ExecuteContext \*context)

## (四) 具体实现

### 1. Parser语法分析

上述功能函数传入的参数均含语法树的结点pSyntaxNode ast，因此，在正式实现函数代码前，要先分析parser语法树的实现。

对于简单的sql语句，如drop table，并不怎么需要分析语法树，因此这里以create index为例。在src/include/parser/minisql.y下找到对应语法，如下：

---

```

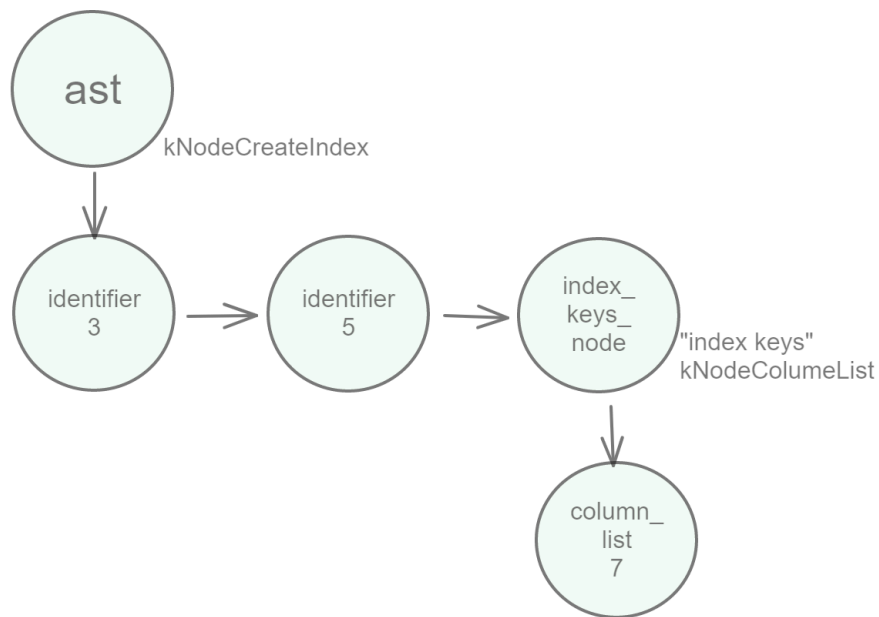
1  sql_create_index:
2      CREATE INDEX IDENTIFIER ON IDENTIFIER '(' column_list ')' {
3          $$ = CreateSyntaxNode(kNodeCreateIndex, NULL);
4          SyntaxNodeAddChildren($$, $3);
5          SyntaxNodeAddChildren($$, $5);
6          pSyntaxNode index_keys_node = CreateSyntaxNode(kNodeColumnList,
7              "index keys");
8          SyntaxNodeAddChildren(index_keys_node, $7);
9          SyntaxNodeAddChildren($$, index_keys_node);
10     }
11 | CREATE INDEX IDENTIFIER ON IDENTIFIER '(' column_list ')' USING
12   IDENTIFIER {
13     $$ = CreateSyntaxNode(kNodeCreateIndex, NULL);
14     SyntaxNodeAddChildren($$, $3);
15     SyntaxNodeAddChildren($$, $5);
16     pSyntaxNode index_keys_node =
17         CreateSyntaxNode(kNodeColumnList, "index keys");
18     SyntaxNodeAddChildren(index_keys_node, $7);
19     SyntaxNodeAddChildren($$, index_keys_node);
20     pSyntaxNode index_type_node =
21         CreateSyntaxNode(kNodeIndexType, "index type");
22     SyntaxNodeAddChildren(index_type_node, $10);
23     SyntaxNodeAddChildren($$, index_type_node);
24 }
25 ;

```

---

不难分析出对应语法树1如下图：





据此可以通过函数传入的语法树结点获取sql语句蕴含的对应信息：

---

```

1  // analyze syntaxNode
2  auto index_name = ast->child->val_; // see parser/minisql.y
3  auto table_name = ast->child->next->val_;
4  auto list = ast->child->next->next_;
  
```

---

## 2. ExecuteCreateTable

CreateTable, DropTable等函数实现较为类似，这里仅以CreateTable函数为例。首先对语法树进行如上分析，获取sql指令中对应的信息。然后根据每一列给出的信息创建表内的属性。具体通过结点状态的不同（kNodeColumnList和kNodeColumnDefinition），先将primary key放入向量表中，再记录其他的key（如若遇到unique key也放入特定向量表中记录）。最后创建表和索引，并将相关信息记入。

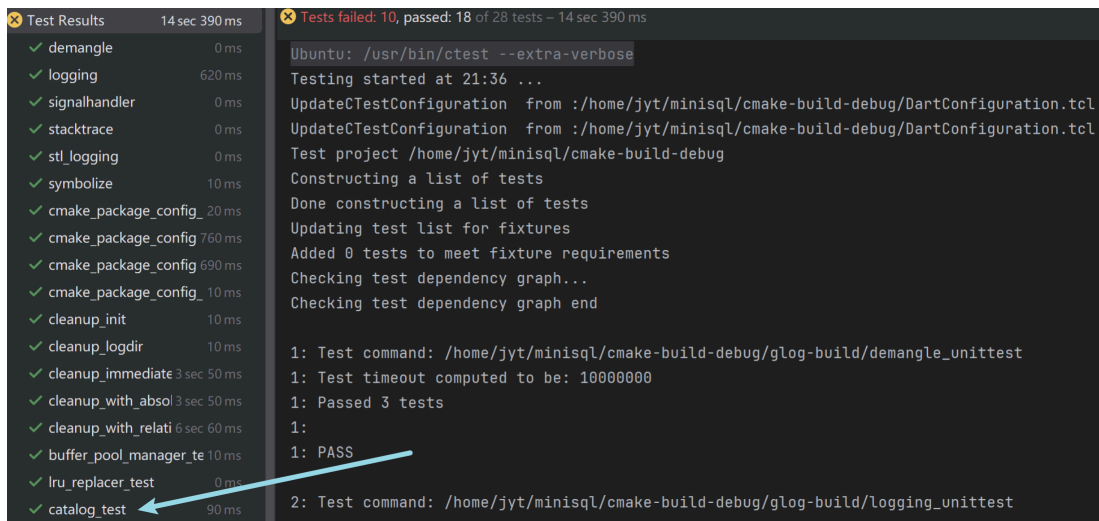
## 3. Executefile

该函数需要读取文件内的sql语句，因此不仅用到了c++的文件流读取函数，还用到了上述提到的语法分析，需要在Parser模块调用yyparse()（一个示例在src/main.cpp中）完成SQL语句解析，得到语法树的根结点pSyntaxNode。

这个函数首先将文件内字符读到缓冲区中，以;为分隔符形成一条或多条sql语句。然后利用parser进行语法解析，执行对应的操作（可以调用给出的Execute函数来进行），最后返回相关信息并关闭文件。

## 四、模块功能验证

### (一) Catalog Manager



The screenshot shows the CMake test results for the 'catalog\_test' module. The test results list shows 'catalog\_test' as passed. The terminal output shows the test command and the result 'PASS'.

```
Test Results 14 sec 390 ms
✓ demangle 0 ms
✓ logging 620 ms
✓ signalhandler 0 ms
✓ stacktrace 0 ms
✓ stl_logging 0 ms
✓ symbolize 10 ms
✓ cmake_package_config_20 ms
✓ cmake_package_config_760 ms
✓ cmake_package_config_690 ms
✓ cmake_package_config_10 ms
✓ cleanup_init 10 ms
✓ cleanup_logdir 10 ms
✓ cleanup_immediate 3 sec 50 ms
✓ cleanup_with_absol 3 sec 50 ms
✓ cleanup_with_relati 6 sec 60 ms
✓ buffer_pool_manager_te 10 ms
✓ lru_replacer_test 0 ms
✓ catalog_test 90 ms

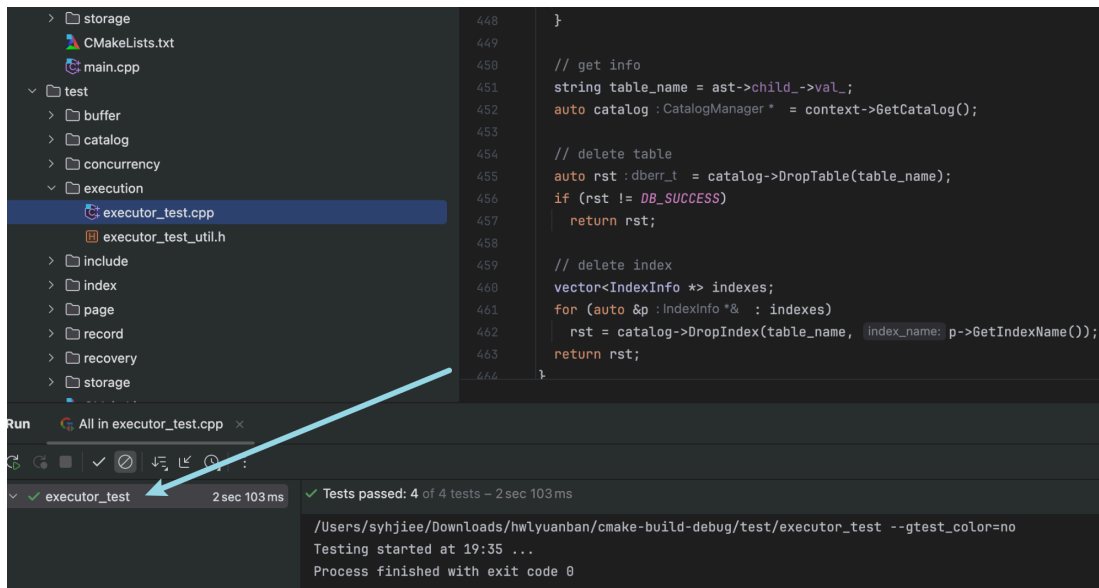
Tests failed: 10, passed: 18 of 28 tests - 14 sec 390 ms

Ubuntu: /usr/bin/ctest --extra-verbose
Testing started at 21:36 ...
UpdateCTestConfiguration from :/home/jyt/minisql/cmake-build-debug/DartConfiguration.tcl
UpdateCTestConfiguration from :/home/jyt/minisql/cmake-build-debug/DartConfiguration.tcl
Test project /home/jyt/minisql/cmake-build-debug
Constructing a list of tests
Done constructing a list of tests
Updating test list for fixtures
Added 0 tests to meet fixture requirements
Checking test dependency graph...
Checking test dependency graph end

1: Test command: /home/jyt/minisql/cmake-build-debug/glog-build/demangle_unittest
1: Test timeout computed to be: 10000000
1: Passed 3 tests
1:
1: PASS

2: Test command: /home/jyt/minisql/cmake-build-debug/glog-build/logging_unittest
```

### (二) Execute Engine



The screenshot shows the code for the 'executor\_test.cpp' file and the test results. The test results show 'executor\_test' as passed. The code shows the implementation of the 'CatalogManager' class.

```
448 }
449
450 // get info
451 string table_name = ast->child_->val_;
452 auto catalog : CatalogManager * = context->GetCatalog();
453
454 // delete table
455 auto rst : dberr_t = catalog->DropTable(table_name);
456 if (rst != DB_SUCCESS)
457     return rst;
458
459 // delete index
460 vector<IndexInfo *> indexes;
461 for (auto &p : IndexInfo *& : indexes)
462     rst = catalog->DropIndex(table_name, p->GetIndexName());
463 return rst;
464 }
```

Run All in executor\_test.cpp x

```
✓ executor_test 2 sec 103 ms
Tests passed: 4 of 4 tests - 2 sec 103 ms
/Users/syhjiej/Downloads/hwlyuanban/cmake-build-debug/test/executor_test --gtest_color=no
Testing started at 19:35 ...
Process finished with exit code 0
```

(说明：在设计实现minisql的过程中，电脑的环境配置出了一点问题，导致原来能构建的minisql现在跑不起来了。询问助教后还是没有找到解决办法，截至提交报告时，电脑环境还是跑不了。因此execute engine的测试是在室友电脑上验证的，显示路径和上面有所不同。可以通过对代码/注释的比对来验证跑的确实是我的execute engine模块。)

然而，在后续进行基本操作时，发现模块五的文件读入部分有问题，会出现“Segmentation fault”报错。截至提交报告时，仍没有发现是哪个地方出了bug。(截图如下，同样是借用了室友的电脑。)

```

antuer@HuWenLuo:/mnt/d/Desktop/Git_project/minisql/Minisql3rd/build/bin$ ./main
minisql > create database db0;
[INFO] Sql syntax parse ok!
minisql > create database db1;
[INFO] Sql syntax parse ok!
minisql > show databases;
[INFO] Sql syntax parse ok!
+-----+
| Database |
+-----+
| db1      |
| db0      |
+-----+
minisql > use db0;
[INFO] Sql syntax parse ok!
Database changed
minisql > create table account(
    id int,
    name char(16),
    balance float,
    primary key(id)
);
[INFO] Sql syntax parse ok!
minisql > show tables;
[INFO] Sql syntax parse ok!
+-----+
| Tables_in_db0 |
+-----+
| account        |
+-----+
minisql > execfile "../../sql_gen/account00.txt";
[INFO] Sql syntax parse ok!
----- Execfile -----
Segmentation fault
antuer@HuWenLuo:/mnt/d/Desktop/Git_project/minisql/Minisql3rd/build/bin$

```

## 五、代码实现

### (一) Catalog Manager

#### 1. index.h

---

```

1 void Init(IndexMetadata *meta_data, TableInfo *table_info,
  BufferPoolManager *buffer_pool_manager) {
2     // Step1: init index metadata and table info
3     // Step2: mapping index key to key schema
4     // Step3: call CreateIndex to create the index
5     meta_data_ = meta_data;
6     const Schema *table_schema = table_info->GetSchema();
7     key_schema_ = Schema::ShallowCopySchema(table_schema, meta_data-
  >GetKeyMapping());
8     index_ = CreateIndex(buffer_pool_manager, "bptree");
9 }

```

---

## 2. index.cpp

```
1  uint32_t IndexMetadata::GetSerializedSize() const {
2      return sizeof(uint32_t)*5 + index_name_.length() +
      key_map_.size()*sizeof(uint32_t);
3  }
```

## 3. catalog.cpp

```
1  /**
2   * TODO: Student Implement FIN OK
3   */
4  uint32_t CatalogMeta::GetSerializedSize() const {
5      return sizeof(uint32_t) * (3 + table_meta_pages_.size() * 2 +
      index_meta_pages_.size() * 2);
6  }
7
8  /**
9   * TODO: Student Implement FIN OK
10  */
11  CatalogManager::CatalogManager(BufferPoolManager
      *buffer_pool_manager, LockManager *lock_manager,
12      LogManager *log_manager, bool init)
13      : buffer_pool_manager_(buffer_pool_manager),
      lock_manager_(lock_manager), log_manager_(log_manager) {
14      if (init) {
15          catalog_meta_ = CatalogMeta::NewInstance();
16          FlushCatalogMetaPage();
17      } else {
18          auto meta_page = buffer_pool_manager_-
      >FetchPage(CATALOG_META_PAGE_ID);
19          catalog_meta_ = CatalogMeta::DeserializeFrom(meta_page-
      >GetData());
20
21          // get info of tables
22          for (auto &p : catalog_meta_->table_meta_pages_) { //
      p.first:table_id; p.second:page_id
23              TableMetadata *table_meta;
24              TableMetadata::DeserializeFrom(buffer_pool_manager_-
      >FetchPage(p.second)->GetData(), table_meta);
25              auto table_heap = TableHeap::Create(buffer_pool_manager_,
      table_meta->GetFirstPageId(),
26
      table_meta->GetSchema(),
      log_manager_, lock_manager_);
27              auto table_info = TableInfo::Create();
28              table_info->Init(table_meta, table_heap);
29              tables_[p.first] = table_info; // store table info
```

```

30     table_names_[table_meta->GetTableName()] = p.first;
31 }
32
33 // get info of indexes
34 for (auto &p : catalog_meta->index_meta_pages_) { //
p.first:index_id; p.second:page_id
35     IndexMetadata *index_meta;
36     IndexMetadata::DeserializeFrom(buffer_pool_manager_-
>FetchPage(p.second)->GetData(), index_meta);
37     auto table_id = index_meta->GetTableId();
38     auto table_info = tables_[table_id];
39     auto table_name = table_info->GetTableName();
40
41     auto index_info = IndexInfo::Create();
42     index_info->Init(index_meta, table_info,
buffer_pool_manager_);
43
44     indexes_[p.first] = index_info; // store index info
45     if (index_names_.find(table_name) == index_names_.end())
46         index_names_.emplace(table_name, unordered_map<string,
index_id_t>());
47     index_names_[table_name][index_meta->GetIndexName()] =
p.first; /**
48     }
49 }
50 }
51
52 /**
53  * TODO: Student Implement FIN OK
54  */
55 dberr_t CatalogManager::CreateTable(const string &table_name,
TableSchema *schema, Txn *txn, TableInfo *&table_info) {
56     if (table_names_.find(table_name) != table_names_.end())
57         return DB_TABLE_ALREADY_EXIST;
58
59     IndexSchema *schema_copy = Schema::DeepCopySchema(schema); /**
debug
60
61     // new id of table and page
62     auto table_id = catalog_meta->GetNextTableId(); // get new id
of table
63     table_names_[table_name] = table_id;
64     page_id_t page_id;
65     auto new_table_page = buffer_pool_manager_->NewPage(page_id); //
get new id of page
66     catalog_meta->table_meta_pages_[table_id] = page_id; // set new
id of table and page
67

```

```

68     // new table info
69     auto table_heap = TableHeap::Create(buffer_pool_manager_,
schema_copy, txn, log_manager_, lock_manager_);
70     auto table_meta = TableMetadata::Create(table_id, table_name,
page_id, schema_copy);
71     table_meta->SerializeTo(new_table_page->GetData());
72     table_info = TableInfo::Create();
73     table_info->Init(table_meta, table_heap);
74     tables_[table_id] = table_info; // set new info of table
75     buffer_pool_manager_->UnpinPage(page_id, true); // unpin page in
buffer
76
77     // update catalog_meta_page
78     auto meta_page = buffer_pool_manager_-
>FetchPage(CATALOG_META_PAGE_ID);
79     catalog_meta_->SerializeTo(meta_page->GetData());
80     buffer_pool_manager_->UnpinPage(CATALOG_META_PAGE_ID, true);
81
82     return DB_SUCCESS;
83 }
84
85 /**
86  * TODO: Student Implement FIN OK
87  */
88 dberr_t CatalogManager::GetTable(const string &table_name,
TableInfo *&table_info) {
89     if (table_names_.find(table_name) == table_names_.end())
90         return DB_TABLE_NOT_EXIST;
91     auto table_id = table_names_[table_name];
92     table_info = tables_[table_id];
93     return DB_SUCCESS;
94 }
95
96 /**
97  * TODO: Student Implement FIN OK
98  */
99 dberr_t CatalogManager::GetTables(vector<TableInfo *> &tables)
const {
100     if (tables_.empty())
101         return DB_TABLE_NOT_EXIST;
102     for(auto &p : tables_)
103         tables.emplace_back(p.second);
104     return DB_SUCCESS;
105 }
106
107 /**
108  * TODO: Student Implement FIN OK
109  */

```

```

110 dberr_t CatalogManager::CreateIndex(const std::string &table_name,
    const string &index_name,
111                                     const std::vector<std::string>
    &index_keys, Txn *txn, IndexInfo *&index_info,
112                                     const string &index_type) {
113     if (table_names_.find(table_name) == table_names_.end())
114         return DB_TABLE_NOT_EXIST;
115     if (index_names_[table_name].find(index_name) !=
    index_names_[table_name].end())
116         return DB_INDEX_ALREADY_EXIST;
117
118     // check whether keys are in the column
119     auto table_id = table_names_[table_name];
120     auto table_info = tables_.at(table_id);
121     auto schema = table_info->GetSchema();
122     std::vector<uint32_t> key_map; // record index_id
123     for (auto &p : index_keys) {
124         index_id_t index_id;
125         if (schema->GetColumnIndex(p, index_id) ==
    DB_COLUMN_NAME_NOT_EXIST)
126             return DB_COLUMN_NAME_NOT_EXIST;
127         key_map.emplace_back(index_id);
128     }
129
130     // new id of index and page
131     auto index_id = catalog_meta_->GetNextIndexId();
132     index_names_[table_name][index_name] = index_id; /* debug .at()
    is wrong
133     page_id_t page_id;
134     auto new_index_page = buffer_pool_manager_->NewPage(page_id); //
    get new id of page
135     catalog_meta_->index_meta_pages_[index_id] = page_id; // set new
    id of index and page
136
137     // new index info
138     auto index_meta = IndexMetadata::Create(index_id, index_name,
    table_id, key_map);
139     index_meta->SerializeTo(new_index_page->GetData());
140     index_info = IndexInfo::Create();
141     index_info->Init(index_meta, table_info, buffer_pool_manager_);
142     indexes_[index_id] = index_info; // set new info of index
143     buffer_pool_manager_->UnpinPage(page_id, true); // unpin page in
    buffer
144
145     // update catalog_meta_page
146     auto meta_page = buffer_pool_manager_-
    >FetchPage(CATALOG_META_PAGE_ID);
147     catalog_meta_->SerializeTo(meta_page->GetData());

```

```

148     buffer_pool_manager_ ->UnpinPage(CATALOG_META_PAGE_ID, true);
149
150     return DB_SUCCESS;
151 }
152
153 /**
154  * TODO: Student Implement FIN OK
155  */
156 dberr_t CatalogManager::GetIndex(const std::string &table_name,
157     const std::string &index_name,
158     IndexInfo *&index_info) const {
159     if (table_names_.find(table_name) == table_names_.end())
160         return DB_TABLE_NOT_EXIST;
161     if (index_names_.at(table_name).find(index_name) ==
162         index_names_.at(table_name).end()) /* [table_name] X
163         return DB_INDEX_NOT_FOUND;
164
165     auto index_id = index_names_.at(table_name).at(index_name);
166     index_info = indexes_.at(index_id);
167     return DB_SUCCESS;
168 }
169
170 /**
171  * TODO: Student Implement FIN OK
172  */
173 dberr_t CatalogManager::GetTableIndexes(const std::string
174     &table_name, std::vector<IndexInfo *> &indexes) const {
175     if (table_names_.find(table_name) == table_names_.end())
176         return DB_TABLE_NOT_EXIST;
177
178     auto index_name = index_names_.at(table_name);
179     for (auto &p : indexes_) {
180         if (table_name == p.second->GetIndexName())
181             indexes.emplace_back(p.second);
182     }
183     return DB_SUCCESS;
184 }
185
186 /**
187  * TODO: Student Implement FIN OK
188  */
189 dberr_t CatalogManager::DropTable(const string &table_name) {
190     if (table_names_.find(table_name) == table_names_.end())
191         return DB_TABLE_NOT_EXIST;
192
193     auto table_id = table_names_[table_name];
194     auto page_id = catalog_meta_ ->table_meta_pages_[table_id];
195     table_names_.erase(table_name); // delete table info

```



```

193     tables_.erase(table_id);
194     catalog_meta_>table_meta_pages_.erase(table_id);
195     buffer_pool_manager_>DeletePage(page_id); // delete page info
196     return DB_SUCCESS;
197 }
198
199 /**
200  * TODO: Student Implement FIN OK
201  */
202 dberr_t CatalogManager::DropIndex(const string &table_name, const
string &index_name) {
203     if (table_names_.find(table_name) == table_names_.end())
204         return DB_TABLE_NOT_EXIST;
205     if (index_names_[table_name].find(index_name) ==
index_names_.at(table_name).end())
206         return DB_INDEX_NOT_FOUND;
207
208     auto index_id = index_names_.at(table_name).at(index_name);
209     auto page_id = catalog_meta_>index_meta_pages_[index_id];
210     index_names_.at(table_name).erase(index_name); // delete index
info
211     indexes_.erase(index_id);
212     catalog_meta_>index_meta_pages_.erase(index_id);
213     buffer_pool_manager_>DeletePage(page_id); // delete page info
214     return DB_SUCCESS;
215 }
216
217 /**
218  * TODO: Student Implement FIN OK
219  */
220 dberr_t CatalogManager::FlushCatalogMetaPage() const {
221     auto page_meta = buffer_pool_manager_
>FetchPage(CATALOG_META_PAGE_ID);
222     catalog_meta_>SerializeTo(page_meta->GetData());
223     if(!buffer_pool_manager_>FetchPage(CATALOG_META_PAGE_ID))
224         return DB_FAILED;
225     return DB_SUCCESS;
226 }
227
228 /**
229  * TODO: Student Implement FIN OK
230  */
231 dberr_t CatalogManager::LoadTable(const table_id_t table_id, const
page_id_t page_id) {
232     if (tables_.find(table_id) != tables_.end())
233         return DB_TABLE_ALREADY_EXIST;
234
235     // get data info from page

```

```

236     catalog_meta_>table_meta_pages_.at(table_id) = page_id; // set
page id
237     TableMetadata *table_meta;
238     TableMetadata::DeserializeFrom(buffer_pool_manager_-
>FetchPage(page_id)->GetData(), table_meta);
239     table_names_.at(table_meta->GetTableName()) = table_id; // store
table name and id
240
241     // load table info
242     auto table_heap = TableHeap::Create(buffer_pool_manager_,
table_meta->GetFirstPageId(),
243                                     table_meta->GetSchema(),
Log_manager_, Lock_manager_);
244     auto table_info = TableInfo::Create();
245     table_info->Init(table_meta, table_heap);
246     tables_[table_id] = table_info; // store table info
247
248     return DB_SUCCESS;
249 }
250
251 /**
252  * TODO: Student Implement FIN OK
253  */
254 dberr_t CatalogManager::LoadIndex(const index_id_t index_id, const
page_id_t page_id) {
255     if (indexes_.find(index_id) != indexes_.end())
256         return DB_INDEX_ALREADY_EXIST;
257
258     // get data info from page
259     catalog_meta_>index_meta_pages_.at(index_id) = page_id; // set
page id
260     IndexMetadata *index_meta;
261     IndexMetadata::DeserializeFrom(buffer_pool_manager_-
>FetchPage(page_id)->GetData(), index_meta);
262     auto table_id = index_meta->GetTableId();
263     auto table_name = tables_[table_id]->GetTableName();
264     index_names_.at(table_name).at(index_meta->GetIndexName()) =
index_id; // store index name and id
265
266     // load index info
267     auto index_info = IndexInfo::Create();
268     index_info->Init(index_meta, tables_[table_id],
buffer_pool_manager_);
269     indexes_[index_id] = index_info; // store index info
270
271     return DB_SUCCESS;
272 }
273

```

```

274  /**
275   * TODO: Student Implement FIN OK
276   */
277  dberr_t CatalogManager::GetTable(const table_id_t table_id,
    TableInfo *&table_info) {
278      if (tables_.find(table_id) == tables_.end())
279          return DB_TABLE_NOT_EXIST;
280      table_info = tables_[table_id];
281      return DB_SUCCESS;
282  }

```

---

## (二) Planner and executor

### 1. execute\_engine.cpp

---

```

1  /**
2   * TODO: Student Implement FIN
3   */
4  dberr_t ExecuteEngine::ExecuteCreateTable(pSyntaxNode ast,
    ExecuteContext *context) {
5      #ifdef ENABLE_EXECUTE_DEBUG
6          LOG(INFO) << "ExecuteCreateTable" << std::endl;
7      #endif
8      if (current_db_.empty()) {
9          cout << "No database selected." << endl;
10         return DB_FAILED;
11     }
12
13     // analyze syntaxNode: parser/minisql.y
14     string table_name = ast->child_->val_;
15     if (ast->child_->type_ != kNodeIdentifier)
16         return DB_FAILED;
17     auto list_head = ast->child_->next_;
18     if (list_head->type_ != kNodeColumnDefinitionList)
19         return DB_FAILED;
20     auto list_node = list_head->child_; // first node now
21
22     vector<string> primary_keys;
23     vector<string> unique_keys;
24     uint32_t index = 0;
25     vector<Column *> columns;
26
27     // get primary key
28     for (; list_node != nullptr; list_node = list_node->next_) {
29         if (list_node->type_ == kNodeColumnList && string(list_node->
>val_) == "primary keys") {
30             for (auto p = list_node->child_; p != nullptr; p = p->next_)
31                 primary_keys.emplace_back(string(p->val_));

```

```

32     }
33 }
34
35 // get other keys
36 for (list_node = list_head->child_; list_node != nullptr;
list_node = list_node->next_) {
37     if (list_node->type_ != kNodeColumnDefinition)
38         continue;
39
40     //- parser analyze
41     bool is_unique = (list_node->val_ != nullptr &&
string(list_node->val_) == "unique");
42     auto deeper_node = list_node->child_;
43     string column_name = deeper_node->val_;
44     string column_type = deeper_node->next_->val_;
45     Column *column;
46
47     //- create column ( defined by column type )
48     if (column_type == "int")
49         column = new Column(column_name, kTypeInt, index++, true,
is_unique);
50     else if (column_type == "float")
51         column = new Column(column_name, kTypeFloat, index++, true,
is_unique);
52     else if (column_type == "char") {
53         string length = deeper_node->next_->child_->val_;
54         for (auto &p : length)
55             if (!isdigit(p))
56                 return DB_FAILED; // exam the number is integer or not
57         if (stoi(length) < 0)
58             return DB_FAILED; // exam the number is positive or not
59         column = new Column(column_name, kTypeChar, stoi(length),
index++, true, is_unique);
60     }
61
62     // add column
63     columns.emplace_back(column);
64     if (is_unique)
65         unique_keys.emplace_back(column_name);
66 }
67
68 // create table
69 auto catalog = context->GetCatalog();
70 auto schema = new Schema(columns);
71 TableInfo* table_info;
72 auto rst = catalog->CreateTable(table_name, schema, context-
>GetTransaction(), table_info);
73 if (rst != DB_SUCCESS)

```

```

74     return rst;
75
76     // create index
77     IndexInfo *index_info;
78     for (auto &p : unique_keys) { // unique key
79         string index_name = "UNIQUE_";
80         index_name += p + "_ON_" + table_name;
81         vector<string> index_keys;
82         index_keys.emplace_back(p);
83         rst = catalog->CreateIndex(table_name, index_name, index_keys,
context->GetTransaction(), index_info, "btree");
84         if (rst != DB_SUCCESS)
85             return rst;
86     }
87     if (primary_keys.size() > 0) {
88         string index_name = "AUTO_CREATED_INDEX_OF_";
89         for (auto &p : primary_keys)
90             index_name += p + "_";
91         index_name += "ON_" + table_name;
92         catalog->CreateIndex(table_name, index_name, primary_keys,
context->GetTransaction(), index_info, "btree");
93     }
94     return rst;
95 }
96
97 /**
98  * TODO: Student Implement FIN
99  */
100 dberr_t ExecuteEngine::ExecuteDropTable(pSyntaxNode ast,
ExecuteContext *context) {
101 #ifdef ENABLE_EXECUTE_DEBUG
102     LOG(INFO) << "ExecuteDropTable" << std::endl;
103 #endif
104     if (current_db.empty()) {
105         cout << "No database selected." << endl;
106         return DB_FAILED;
107     }
108
109     // get info
110     string table_name = ast->child_->val_;
111     auto catalog = context->GetCatalog();
112
113     // delete table
114     auto rst = catalog->DropTable(table_name);
115     if (rst != DB_SUCCESS)
116         return rst;
117
118     // delete index

```

```

119     vector<IndexInfo *> indexes;
120     for (auto &p : indexes)
121         rst = catalog->DropIndex(table_name, p->GetIndexName());
122     return rst;
123 }
124
125 /**
126  * TODO: Student Implement FIN
127  */
128 dberr_t ExecuteEngine::ExecuteShowIndexes(pSyntaxNode ast,
129 ExecuteContext *context) {
130 #ifdef ENABLE_EXECUTE_DEBUG
131     LOG(INFO) << "ExecuteShowIndexes" << std::endl;
132 #endif
133     if (current_db_.empty()) {
134         cout << "No database selected." << endl;
135         return DB_FAILED;
136     }
137
138     // get tables' info
139     auto catalog = context->GetCatalog();
140     vector<TableInfo *> tables;
141     catalog->GetTables(tables);
142
143     // show indexes
144     cout << "----- Show Index -----" << endl;
145     int counter = 1;
146     for (auto &p : tables) {
147         vector<IndexInfo *> indexes;
148         catalog->GetTableIndexes(p->GetTableName(), indexes);
149         cout << "\tTable: " << p->GetTableName() << endl;
150         for (auto &i : indexes) {
151             cout << "\t\tIndex " << (counter++) << ": " << i-
152 >GetIndexName() << endl;
153         }
154     }
155     cout << "Totally" << counter << "index(es) listed." << endl;
156     return DB_SUCCESS;
157 }
158
159 /**
160  * TODO: Student Implement FIN
161  */
162 dberr_t ExecuteEngine::ExecuteCreateIndex(pSyntaxNode ast,
163 ExecuteContext *context) {
164 #ifdef ENABLE_EXECUTE_DEBUG
165     LOG(INFO) << "ExecuteCreateIndex" << std::endl;
166 #endif

```

```

164     if (current_db_.empty()) {
165         cout << "No database selected." << endl;
166         return DB_FAILED;
167     }
168
169     // analyze syntaxNode
170     auto index_name = ast->child_->val_; // see parser/minisql.y
171     auto table_name = ast->child_->next_->val_;
172     auto list = ast->child_->next_->next_;
173     if (list->type_ != kNodeColumnList)
174         return DB_FAILED;
175
176     // get index info and create index
177     auto catalog = context->GetCatalog();
178     vector<string> keys;
179     for (auto key = list->child_; key != nullptr; key = key->next_)
180     {
181         keys.emplace_back(key->val_);
182     }
183     IndexInfo *index_info;
184     auto rst = catalog->CreateIndex(table_name, index_name, keys,
185     context->GetTransaction(), index_info, "btree");
186     return rst;
187 }
188
189 /**
190  * TODO: Student Implement FIN
191  */
192 dberr_t ExecuteEngine::ExecuteDropIndex(pSyntaxNode ast,
193 ExecuteContext *context) {
194 #ifdef ENABLE_EXECUTE_DEBUG
195     LOG(INFO) << "ExecuteDropIndex" << std::endl;
196 #endif
197     if (current_db_.empty()) {
198         cout << "No database selected." << endl;
199         return DB_FAILED;
200     }
201
202     // get info
203     string index_name = ast->child_->val_;
204     auto catalog = context->GetCatalog();
205     vector<TableInfo *> tables;
206     catalog->GetTables(tables);
207
208     // drop index
209     for (auto &p : tables)
210         if (catalog->DropIndex(p->GetTableName(), index_name) !=
211 DB_SUCCESS)

```

```

208         return DB_FAILED;
209     return DB_SUCCESS;
210 }
211
212 dberr_t ExecuteEngine::ExecuteTrxBegin(pSyntaxNode ast,
ExecuteContext *context) {
213 #ifdef ENABLE_EXECUTE_DEBUG
214     LOG(INFO) << "ExecuteTrxBegin" << std::endl;
215 #endif
216     return DB_FAILED;
217 }
218
219 dberr_t ExecuteEngine::ExecuteTrxCommit(pSyntaxNode ast,
ExecuteContext *context) {
220 #ifdef ENABLE_EXECUTE_DEBUG
221     LOG(INFO) << "ExecuteTrxCommit" << std::endl;
222 #endif
223     return DB_FAILED;
224 }
225
226 dberr_t ExecuteEngine::ExecuteTrxRollback(pSyntaxNode ast,
ExecuteContext *context) {
227 #ifdef ENABLE_EXECUTE_DEBUG
228     LOG(INFO) << "ExecuteTrxRollback" << std::endl;
229 #endif
230     return DB_FAILED;
231 }
232
233 /**
234  * TODO: Student Implement FIN
235  */
236
237 dberr_t ExecuteEngine::ExecuteExecfile(pSyntaxNode ast,
ExecuteContext *context) {
238 #ifdef ENABLE_EXECUTE_DEBUG
239     LOG(INFO) << "ExecuteExecfile" << std::endl;
240 #endif
241     auto file_name = ast->child_->val_;
242     fstream file(file_name);
243     if (!file.is_open()) {
244         cout << "Fail to open file '" << file_name << "'. " << endl;
245         return DB_FAILED;
246     }
247
248     // buffer for cmd
249     const int buffer_size = 1024;
250     char *cmd = new char[buffer_size];
251

```



```

252     cout << "----- Execfile -----" << endl;
253     auto start_time = std::chrono::system_clock::now(); /**
254     while (!file.eof()) {
255         // load cmd
256         memset(cmd, 0, buffer_size);
257         int counter = 0;
258         char c;
259         while (!file.eof() && (c == file.get()) != ';')
260             cmd[counter++] = c;
261         if (file.eof()) break; // has scanned all file
262         cmd[counter] = ';'; // end of one sql
263
264         // translate sql via parser
265         YY_BUFFER_STATE bp = yy_scan_string(cmd);
266         MinisqlParserInit();
267         yyparse();
268         if (MinisqlParserGetError())
269             printf("%s\n", MinisqlParserGetErrorMessage()); // error
message
270
271         // handle result
272         auto rst = Execute(MinisqlGetParserRootNode());
273
274         // parser over, clean memory
275         MinisqlParserFinish();
276         yy_delete_buffer(bp);
277         yylex_destroy();
278
279         // handle execute message
280         ExecuteInformation(rst);
281     }
282     auto stop_time = std::chrono::system_clock::now();
283
284     // calculate execution time
285     auto time =
double((std::chrono::duration_cast<std::chrono::milliseconds>
(stop_time - start_time)).count());
286     cout << "Execfile finished in " << time << "ms." << endl;
287     return DB_SUCCESS;
288 }
289
290 /**
291  * TODO: Student Implement FIN
292  */
293 dberr_t ExecuteEngine::ExecuteQuit(pSyntaxNode ast, ExecuteContext
*context) {
294 #ifdef ENABLE_EXECUTE_DEBUG
295     LOG(INFO) << "ExecuteQuit" << std::endl;

```

```
296 #endif
297     current_db_ = "";
298     cout << "Bye." << endl;
299     return DB_SUCCESS;
300 }
```

---